

Heap Heap Hooray: Improving Memory Management

Tyler Gutowski, Trevor Schiff,
Dr. Ryan Stansifer (client)

Task	Tyler	Trevor
Analysis of the Garbage Collector Handbook	0.5	0.5
Evaluation of pros & cons	1.0	0.0
Examination of open-source projects	0.0	1.0
Defining objectives and scope	0.2	0.8
Determining vital metrics	1.0	0.0
Assessment of project feasibility	0.5	0.5
Identification of prospective challenges	0.0	1.0
Selection of a suitable garbage collection algorithm	0.8	0.2
Development of high-level design	0.8	0.2
Development of detailed design documents	0.0	1.0

Algorithm	Pro	Cons
Reference Counting	Simple Real-time reclamation	Counter costs memory Stack variables are inefficient
Mark-And-Sweep	GC operations only occur in GC cycles	All memory must be searched during cycles Memory fragmented
Copying	Defragments memory	Only a portion of memory is usable
Generational	Young is collected frequently Fast GC cycles	Tuning required Only a portion of memory is usable
Incremental	No stutters	GC is slower
Concurrent/Parallel	No GC cycle stutters	Impractical

Basic Overview: Reference Counting

Assume objects are initialized with a reference count of 0

If a reference to an object is created

The object's reference count will increment

If a reference to an object is destroyed

The object's reference count will decrement

If the object's reference count is now zero

Memory allocated for the object is now freed

Basic Overview: Mark-and-Sweep

Assume each object initialized has a “mark” bit

If an attempt to allocate a new object fails

- For each root object (stack locals, registers)

- For each child heap object

- Mark the object

- For each object on the heap

- If the object isn't marked

- Free the object

Vital Metrics

- Throughput
 - Allocations per second
 - Reclaims per second
- Pause Times
 - Maximum pause time
 - Average pause time
 - Pause time frequency
- Memory Overhead
 - Total memory allocated
 - Total memory available
- CPU Overhead
 - Resources used by application
 - Resources used by garbage collector

Feasibility

- Milestone 1 completely completed
- Milestone 2 is definitely feasible
- Make sure to realistically plan for Milestone 2
- Primary and secondary goals
 1. Reference Counting
 2. Mark-and-Sweep
- Wait until progress on Milestone 2 to begin planning Milestone 3

Design

- Main focus on reference counting
 - Tally all live references to a given block
 - Free block when ref count is decremented to zero (unreachable)
- Mark-sweep GC as secondary implementation
 - RC cannot collect cyclic (self-referencing) garbage
 - Traverses heap block graph by searching block data for pointers
- Implemented in MiniJava runtime, in C
 - Compiler inserts calls during codegen phase
- Header created at start of every heap-allocated block
 - For reference counting: ref count
 - For mark-sweep: size, next block, mark flag

Heap Block

```
// 32-bit "tag" to identify heap blocks from all other memory
```

```
#define HEAP_BLOCK_TAG 'HBLK'
```

```
typedef struct HeapBlock {  
    // Block identification  
    u32 tag;                // at 0x0  
  
    // Next block in list of runtime allocations  
    struct HeapBlock* next; // at 0x4  
    // Size of this allocation  
    size_t size;            // at 0x8  
  
    // Mark bit (for mark-sweep GC)  
    s32 marked : 1;         // at 0xC  
    // Reference count (for reference count GC)  
    s32 ref : 31;           // at 0xC  
  
    // Block user data begins at offset 0x10 . . .  
} HeapBlock;
```

Milestone 2 Plans

- Compiler integration
 - Generate calls to runtime GC functions
- Runtime integration
 - Implement reference counting
 - Increment, decrement functions
 - Begin mark-sweep implementation if time allows
 - Mark, sweep functions
- Performance analysis